

```

import numpy as np
import time

import os

import chainer
from chainer import cuda, Function, gradient_check, Variable,
optimizers, serializers, utils
from chainer import Link, Chain, ChainList
import chainer.functions as F
import chainer.links as L

#load CIFAR10 data
import h5py
CIFAR10_data = h5py.File('CIFAR10.hdf5', 'r')
x_train = np.float32(CIFAR10_data['X_train'][:])
y_train = np.int32(np.array(CIFAR10_data['Y_train'][:]))
x_test = np.float32(CIFAR10_data['X_test'][:])
y_test = np.int32( np.array(CIFAR10_data['Y_test'][:]) )

CIFAR10_data.close()

#D = 32
num_outputs = 10

class Conv_NN(Chain):
    def __init__(self):
        super(Conv_NN, self).__init__(
            conv1_1=L.Convolution2D(3, 32, 3, pad=1),
            bn1_1=L.BatchNormalization(32),
            conv1_2=L.Convolution2D(32, 32, 3, pad=1),
            bn1_2=L.BatchNormalization(64),
            conv1_3=L.Convolution2D(32, 32, 3, pad=1),
            bn1_3=L.BatchNormalization(64),
            conv1_4=L.Convolution2D(32, 32, 3, pad=1),
            bn1_4=L.BatchNormalization(32),

            conv2_1=L.Convolution2D(32, 64, 3, pad=1),
            bn2_1=L.BatchNormalization(64),
            conv2_2=L.Convolution2D(64, 64, 3, pad=1),
            bn2_2=L.BatchNormalization(64),

            bn3_1=L.BatchNormalization(128),
            conv3_2=L.Convolution2D(128, 128, 3, pad=1),
            bn3_2=L.BatchNormalization(128),
            conv3_3=L.Convolution2D(128, 128, 3, pad=1),
            bn3_3=L.BatchNormalization(128),
            conv3_4=L.Convolution2D(128, 128, 3, pad=1),
            bn3_3=L.BatchNormalization(128),

```

```

        conv3_5=L.Convolution2D(128, 128, 3, pad=1),
        bn3_5=L.BatchNormization(128),
        conv3_6=L.Convolution2D(128, 128, 3, pad=1),
        bn3_6=L.BatchNormization(128),

        fc4 = L.Linear(2048, 500),
        fc5 = L.Linear(500, 500),
        fc6 = L.Linear(500,10),
    )
def __call__(self, x_data, y_data, dropout_bool, bn_bool, p):
    x = Variable(x_data)
    t = Variable(y_data)
    h = F.relu(self.bn1_1(self.conv1_1(x), bn_bool))
    h = F.relu(self.bn1_2(self.conv1_2(h), bn_bool))
    h = F.relu(self.bn1_3(self.conv1_3(h), bn_bool))
    h = F.relu(self.bn1_4(self.conv1_4(h), bn_bool))
    h = F.max_pooling_2d(h, 3, 2)
    h = F.dropout(h, p, dropout_bool)

    h = F.relu(self.bn2_1(self.conv2_1(h), bn_bool))
    h = F.relu(self.bn2_2(self.conv2_2(h), bn_bool))
    h = F.max_pooling_2d(h, 3, 2)
    h = F.dropout(h, p, dropout_bool)

    h = F.relu(self.bn3_1(self.conv3_1(h), bn_bool))
    h = F.relu(self.bn3_2(self.conv3_2(h), bn_bool))
    h = F.relu(self.bn3_3(self.conv3_3(h), bn_bool))
    h = F.relu(self.bn3_4(self.conv3_4(h), bn_bool))
    h = F.relu(self.bn3_5(self.conv3_5(h), bn_bool))
    h = F.relu(self.bn3_6(self.conv3_6(h), bn_bool))
    h = F.max_pooling_2d(h, 3, 2)
    h = F.dropout(h, p, dropout_bool)

    h = F.dropout(F.relu(self.fc4(h)), p, dropout_bool)
    h = F.dropout(F.relu(self.fc5(h)), p, dropout_bool)
    h = self.fc6(h)
    L_out = h
    return F.softmax_cross_entropy(L_out, t), F.accuracy(L_out, t)

```

#returns test accuracy of the model. dropout is set to its test state

```

def Calculate_Test_Accuracy(x_test, y_test, model, p, GPU_on,
batch_size):
    L_Y_test = len(y_test)
    counter = 0
    test_accuracy_total = 0.0
    for i in range(0, L_Y_test, batch_size):
        if (GPU_on):
            x_batch = cuda.to_gpu(x_test[i:i+ batch_size,:])
            y_batch = cuda.to_gpu(y_test[i:i+ batch_size] )

```

```

        else:
            x_batch = x_test[i:i+batch_size,:]
            y_batch = y_test[i:i+batch_size]
            dropout_bool = False
            bn_bool = True
            loss, accuracy = model(x_batch, y_batch, dropout_bool, bn_bool,
p)
            test_accuracy_batch = 100.0*np.float(accuracy.data )
            test_accuracy_total += test_accuracy_batch
            counter += 1
        test_accuracy = test_accuracy_total/(np.float(counter))
        return test_accuracy

```

```

model = Conv_NN()

```

```

#True if training with GPU, False if training with CPU

```

```

GPU_on = True

```

```

#GPU_on = False

```

```

#size of minibatches

```

```

batch_size = 250

```

```

#transfer model to GPU

```

```

if (GPU_on):

```

```

    model.to_gpu()

```

```

#optimization method

```

```

optimizer = optimizers.MomentumSGD(momentum = .99)

```

```

#optimizer = optimizers.RMSprop(lr=0.001, alpha=0.99, eps=1e-08)

```

```

optimizer.setup(model)

```

```

#learning rate

```

```

optimizer.lr = .01

```

```

#dropout probability

```

```

p = .25

```

```

#number of training epochs

```

```

num_epochs = 300

```

```

L_Y_train = len(y_train)

```

```

time1 = time.time()

```

```

for epoch in range(num_epochs):

```

```

    #reshuffle dataset

```

```

    I_permutation = np.random.permutation(L_Y_train)

```

```

    x_train = x_train[I_permutation,:]

```

```

    y_train = y_train[I_permutation]

```

```

epoch_accuracy = 0.0
batch_counter = 0
for i in range(0, L_Y_train, batch_size):
    if (GPU_on):
        x_batch = cuda.to_gpu(x_train[i:i+batch_size,:])
        y_batch = cuda.to_gpu(y_train[i:i+batch_size] )
    else:
        x_batch = x_train[i:i+batch_size,:]
        y_batch = y_train[i:i+batch_size]
    model.zerograds()
    dropout_bool = True
    bn_bool = False
    loss, accuracy = model(x_batch, y_batch, dropout_bool,
bn_bool, p)
    loss.backward()
    optimizer.update()
    #print("success")
    epoch_accuracy += np.float(accuracy.data)
    batch_counter += 1
    if (epoch % 1 == 0):
        #print "Epoch %d" % epoch
        train_accuracy = 100.0*epoch_accuracy/np.float(batch_counter)
        print "Train accuracy: %f" % train_accuracy
    if (epoch == num_epochs-1):
        test_accuracy = Calculate_Test_Accuracy(x_test, y_test, model,
p, GPU_on, batch_size)
        print "Test Accuracy: %f" % test_accuracy

time2 = time.time()
training_time = time2 - time1
#print "Rank: %d" % rank
print "Training time: %f" % training_time

```